

Socket Programming

Yunmin Go

School of CSEE

Handong Global University

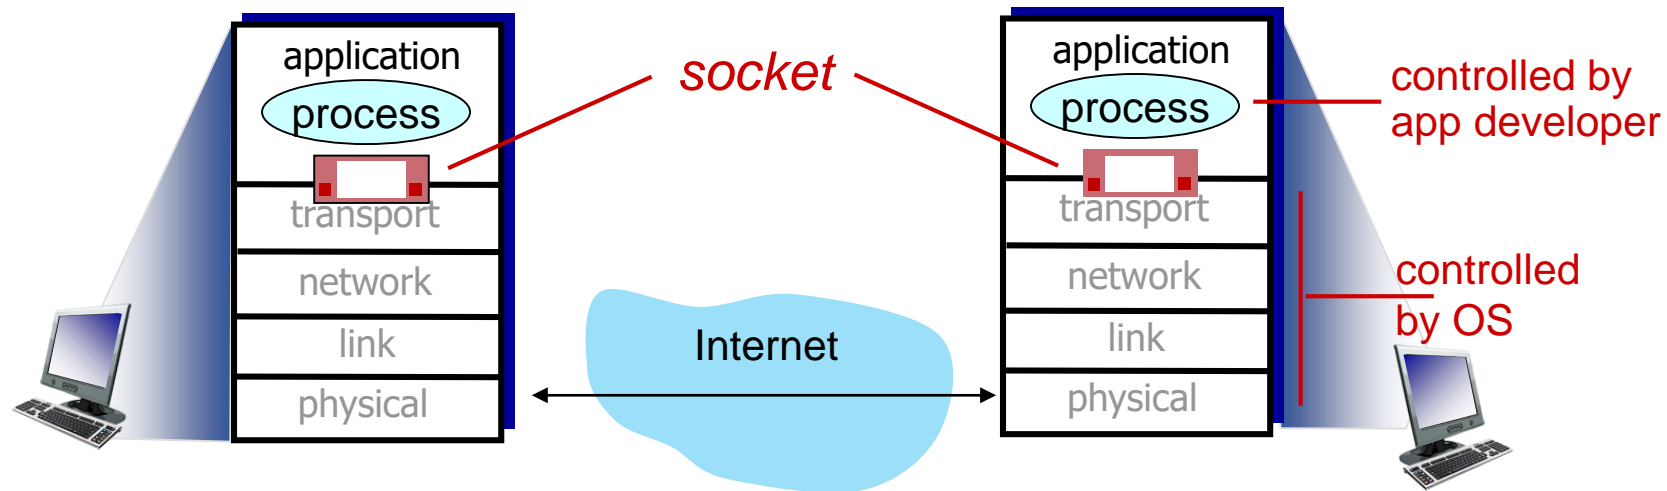
Socket Programming: Roadmap

- Introduction
- TCP Socket programming #1
- TCP Socket programming #2
- UDP Socket programming
- Web Server Example

Socket programming

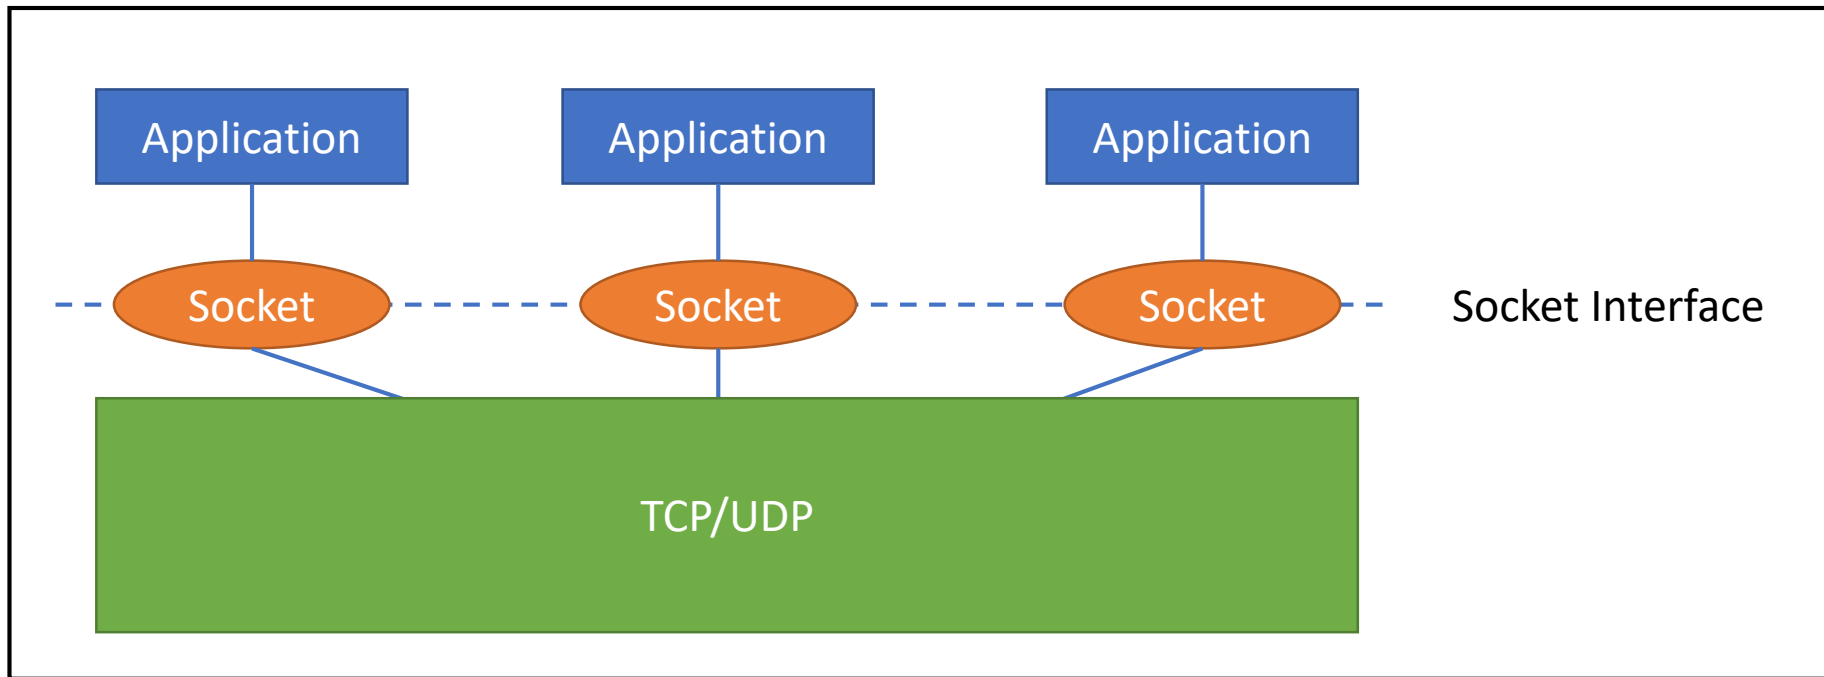
goal: learn how to build client/server applications that communicate using sockets

socket: door between application process and end-end-transport protocol



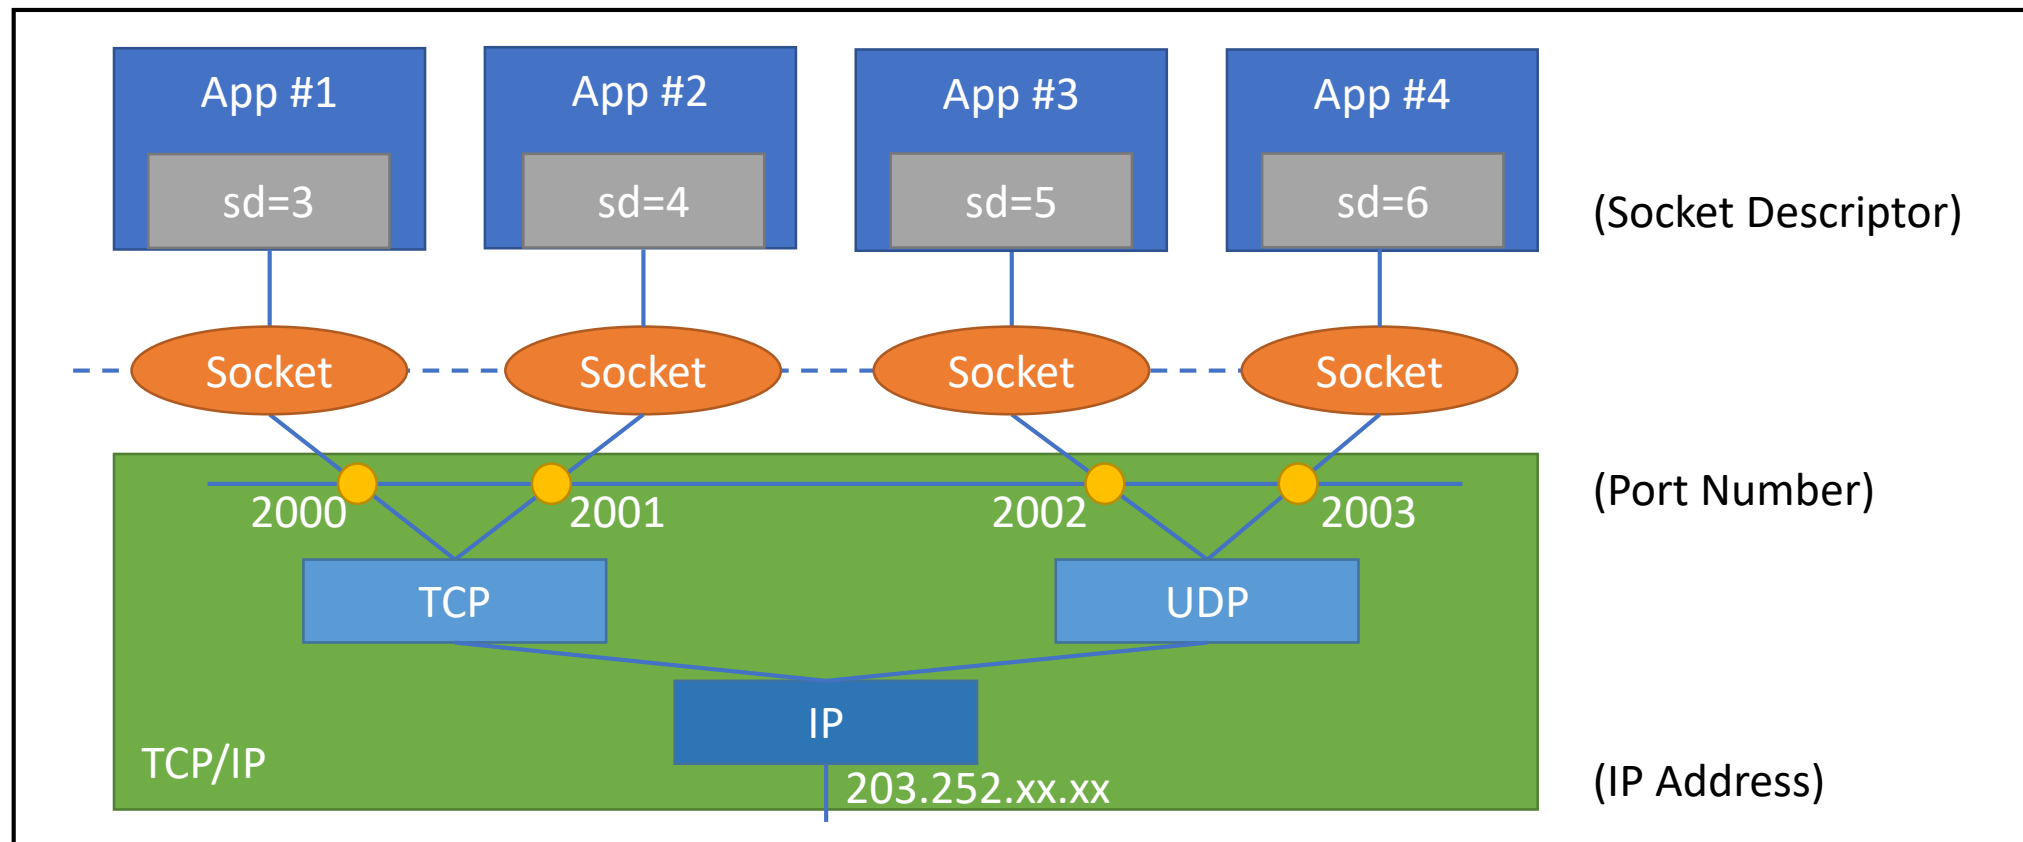
Socket

- Socket interface
 - Located between application and TCP, UDP



Socket

- Socket and TCP/UDP relationship



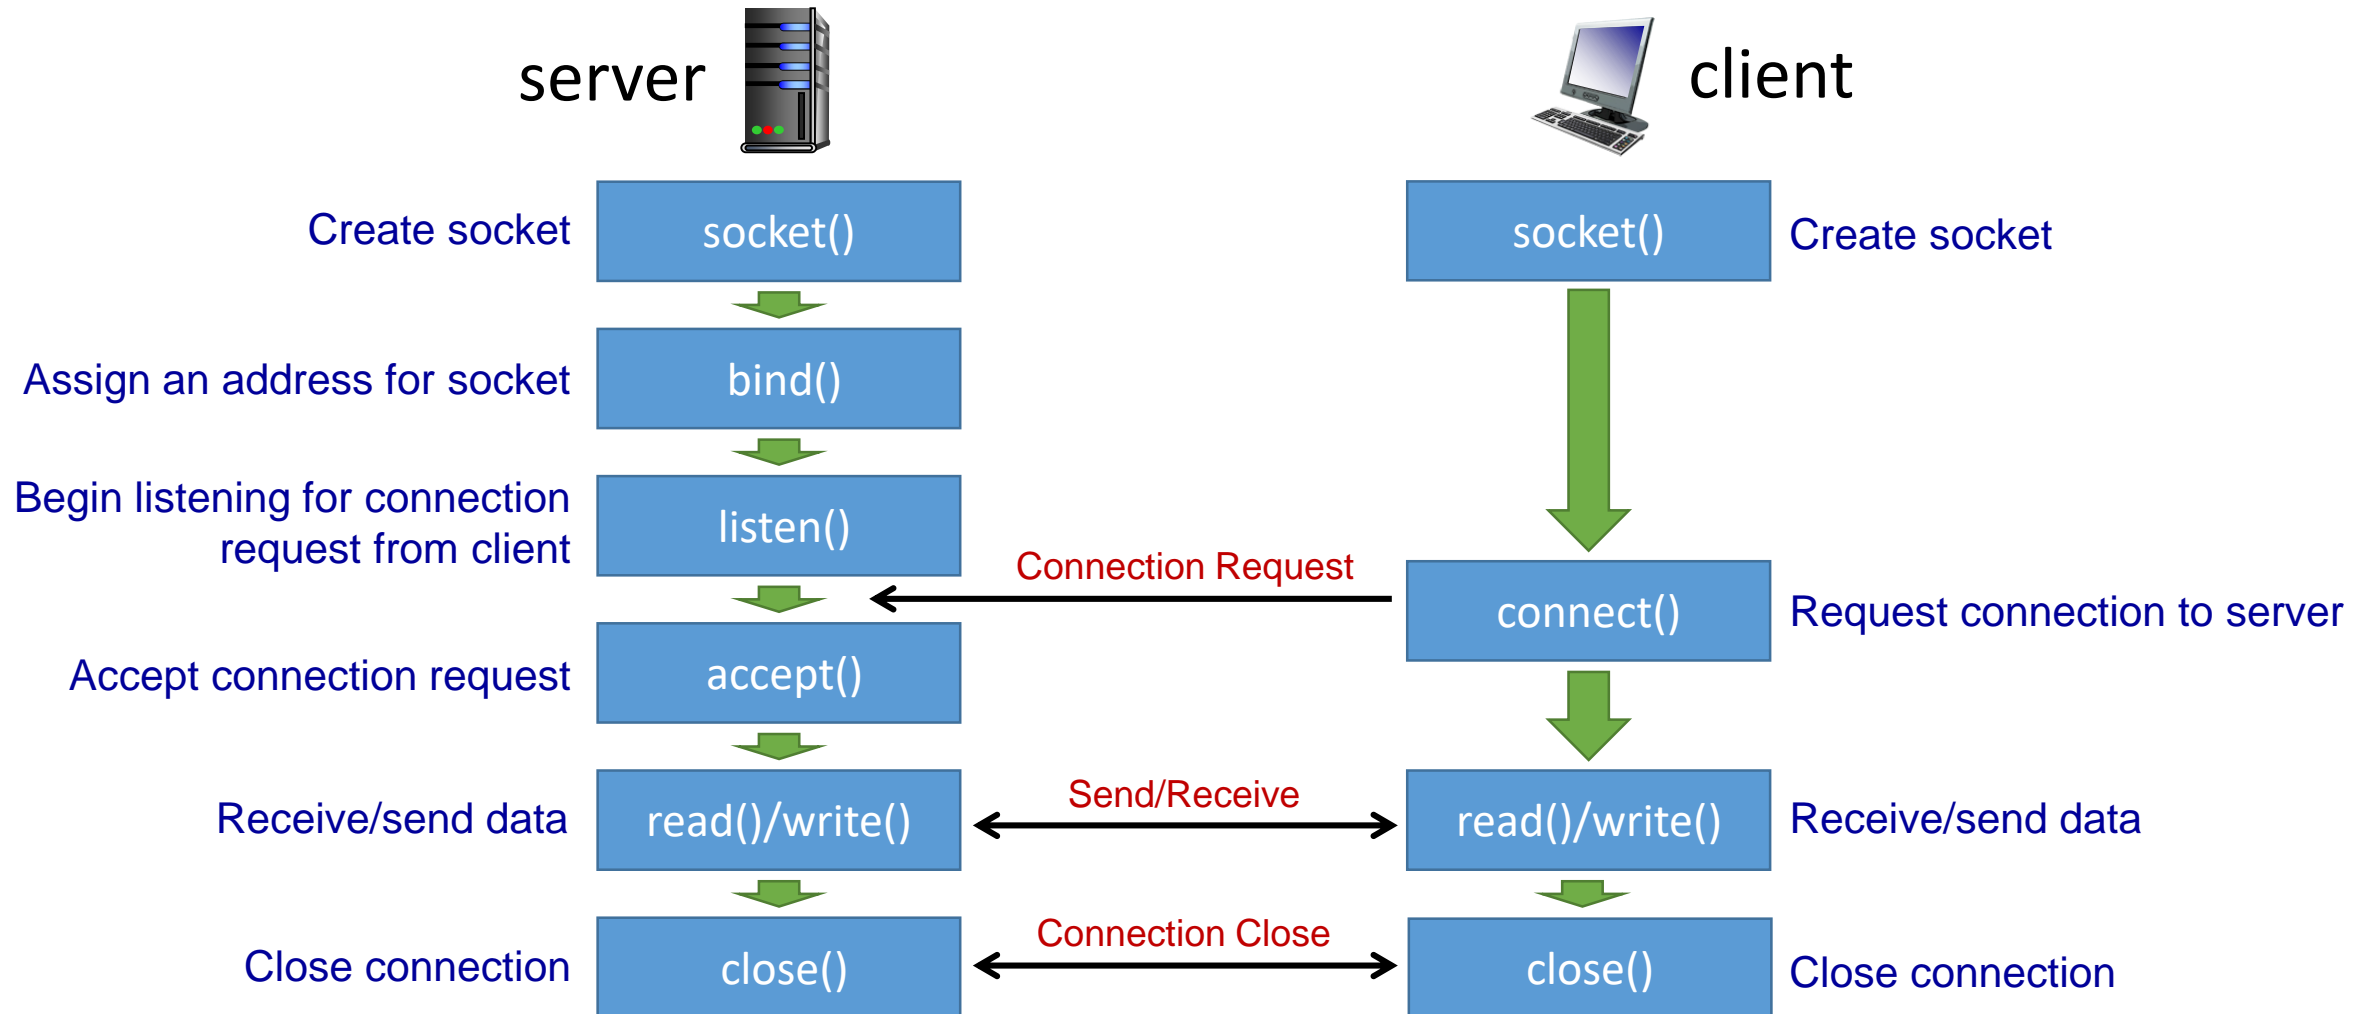
Socket

- Each socket is associated with five components
 - Protocol
 - Protocol family and protocol
 - Source address, source port
 - Destination address, destination port
- Where to define components
 - Protocol: `socket()`
 - Source address and port: `bind()`
 - Destination address and port: `connect()`, `sendto()`

Socket Programming: Roadmap

- Introduction
- TCP Socket programming #1
- TCP Socket programming #2
- UDP Socket programming
- Web Server Example

TCP Server/Client Function Call



TCP Server Example: tcp_server.c

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <arpa/inet.h>
#include <sys/types.h>
#include <sys/socket.h>

void error_handling(char *message);

int main(int argc, char *argv[])
{
    int serv_sock;
    int clnt_sock;

    struct sockaddr_in serv_addr;
    struct sockaddr_in clnt_addr;
    socklen_t clnt_addr_size;

    char message[]="Hello World!";

    if (argc != 2){
        printf("Usage : %s <port>\n", argv[0]);
        exit(1);
    }
```

```
serv_sock = socket(PF_INET, SOCK_STREAM, 0);
if (serv_sock == -1)
    error_handling("socket() error");

memset(&serv_addr, 0, sizeof(serv_addr));
serv_addr.sin_family = AF_INET;
serv_addr.sin_addr.s_addr = htonl(INADDR_ANY);
serv_addr.sin_port = htons(atoi(argv[1]));

if (bind(serv_sock, (struct sockaddr*)&serv_addr, sizeof(serv_addr)) == -1)
    error_handling("bind() error");

if (listen(serv_sock, 5) == -1)
    error_handling("listen() error");

clnt_addr_size = sizeof(clnt_addr);
clnt_sock = accept(serv_sock, (struct sockaddr*)&clnt_addr, &clnt_addr_size);
if (clnt_sock == -1)
    error_handling("accept() error");

write(clnt_sock, message, sizeof(message));

close(clnt_sock);
close(serv_sock);
return 0;
}
```

TCP Client Example: tcp_client.c

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <arpa/inet.h>
#include <sys/types.h>
#include <sys/socket.h>

void error_handling(char *message);

int main(int argc, char* argv[])
{
    int sock;
    struct sockaddr_in serv_addr;
    char message[30];
    int str_len = 0;
    int idx = 0, read_len = 0;

    if (argc != 3) {
        printf("Usage : %s <IP> <port>\n", argv[0]);
        exit(1);
    }
```

```
sock = socket(PF_INET, SOCK_STREAM, 0);
if (sock == -1)
    error_handling("socket() error");
```

socket()

```
memset(&serv_addr, 0, sizeof(serv_addr));
serv_addr.sin_family = AF_INET;
serv_addr.sin_addr.s_addr = inet_addr(argv[1]);
serv_addr.sin_port = htons(atoi(argv[2]));
```

connect()

```
if (connect(sock, (struct sockaddr*)&serv_addr, sizeof(serv_addr)) == -1)
    error_handling("connect() error!");
```

```
while (read_len = read(sock, &message[idx++], 1))
{
    if (read_len == -1)
    {
        error_handling("read() error!");
        break;
    }
    str_len += read_len;
}
```

read()

```
printf("Message from server: %s \n", message);
printf("Function read call count: %d \n", str_len);
```

```
close(sock);
return 0;
```

close()

```
}
```

socket()

```
#include <sys/socket.h>
int socket(int domain, int type, int protocol)
```

Create an endpoint for communication

- *domain*: protocol family
 - PF_INET: IPv4
 - PF_INET6: IPv6
- *type*: type of service
 - SOCK_STREAM: TCP
 - SOCK_DGRAM: UDP
 - SOCK_RAW: raw IP
- *protocol*: specifies the specific protocol
 - Usually 0 which means the default
 - IPPROTO_TCP
 - IPPROTO_UDP
- Return value
 - Success: a file descriptor for the new socket
 - Error: -1
- Example

```
sock = socket(PF_INET, SOCK_STREAM, 0);
if (sock == -1)
    error_handling("socket() error");
```

TCP Socket

```
sock = socket(PF_INET, SOCK_DGRAM, 0);
if (sock == -1)
    error_handling("socket() error");
```

UDP Socket

bind()

```
#include <sys/socket.h>
```

```
int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

Assign an address to a socket

- *sockfd*: file descriptor for the socket
- *addr*: address (IP address and port number) to assign a socket
- *addrlen*: the size (bytes) of the address structure pointed to by *addr*

Return value

- Success: 0
- Error: -1

Example

```
memset(&serv_addr, 0, sizeof(serv_addr));  
serv_addr.sin_family = AF_INET;  
serv_addr.sin_addr.s_addr = htonl(INADDR_ANY);  
serv_addr.sin_port = htons(atoi(argv[1]));  
  
if (bind(serv_sock, (struct sockaddr*) &serv_addr, sizeof(serv_addr)) == -1)  
    error_handling("bind() error");
```

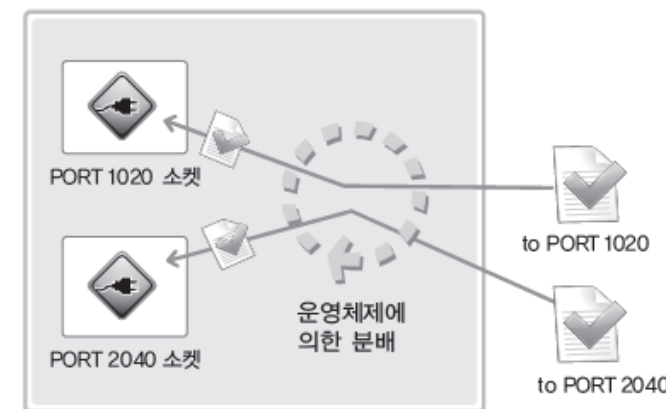
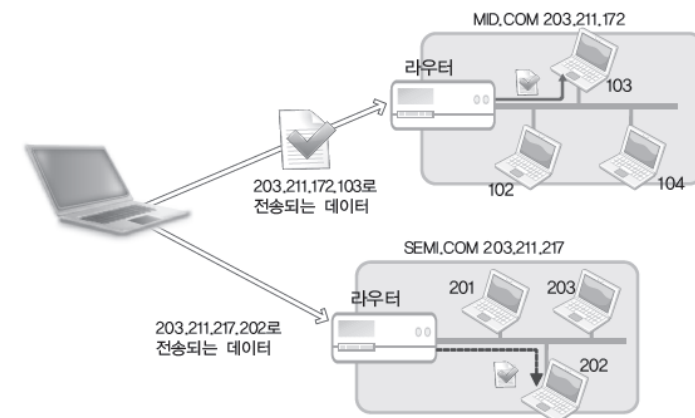
Internet Address

■ IP Address

- Used to identify computers on the Internet
- There are **IPv4** (4-byte address) and IPv6 (16bytes address)
- When creating a socket, we must specify a basic protocol

■ Port

- Used to identify sockets in the computer
- Port number is represented in 16bits
→ Total range 0 ~ 65535
- **Well-known port**: 0 ~ 1023
 - its use has already been decided
 - e.g., 22 for ssh, 80 for web, etc
- **Ephemeral port**: a short-lived transport protocol port
 - allocated automatically from a predefined range by the IP stack software
 - e.g., Linux: 32768 to 60999



Address Structure

- Defined in <netinet/in.h>

```
struct sockaddr {  
    u_short      sa_family;    // Address family  
    char         sa_data[14];  // address  
};
```

```
struct sockaddr_in {  
    sa_family_t   sin_family;   // Address family: AF_INET  
    uint16_t      sin_port;     // Port number (16bits), Network byte order  
    struct in_addr sin_addr;     // IP address (32bits), Network byte order  
    char          sin_zero[8];  // Unused  
};
```

```
struct in_addr {  
    in_addr_t     s_addr;       // IPv4 Internet address (32bits)  
};
```




```
memset(&serv_addr, 0, sizeof(serv_addr));  
serv_addr.sin_family = AF_INET;  
serv_addr.sin_addr.s_addr = htonl(INADDR_ANY);  
serv_addr.sin_port = htons(atoi(argv[1]));  
  
if (bind(serv_sock, (struct sockaddr*) &serv_addr, sizeof(serv_addr)) == -1)  
    error_handling("bind() error");
```

※ INADDR_ANY: the socket will be bound to all local interfaces

Byte Order Conversion

- Byte ordering
 - Little endian: least significant byte first
 - Big endian: most significant byte first
- Network byte order = big endian
 - Most-significant byte at least address of a word
 - Host byte order depends on host's CPU

- 32-bit signed or unsigned integer comprises 4 bytes



htons(), ntohs(), htonl(), ntohl()

```
#include <arpa/inet.h>
uint16_t htons(uint16_t hostshort); // convert unsigned short integer from host byte order to network byte order
uint16_t ntohs(uint16_t netshort); // convert unsigned short integer from network byte order to host byte order
uint32_t htonl(uint32_t hostlong); // convert unsigned integer from host byte to network byte order
uint32_t ntohl(uint32_t netlong); // convert unsigned integer from network byte order to host byte order
```

Convert values between host and network byte order

■ Example

```
unsigned short host_port = 0x1234;
unsigned short net_port;
unsigned long host_addr = 0x12345678;
unsigned long net_addr;

net_port = htons(host_port);
net_addr = htonl(host_addr);

printf("Host ordered port: %#x \n", host_port);
printf("Network ordered port: %#x \n", net_port);
printf("Host ordered address: %#lx \n", host_addr);
printf("Network ordered address: %#lx \n", net_addr);
```

```
>> Host ordered port: 0x1234
>> Network ordered port: 0x3412
>> Host ordered address: 0x12345678
>> Network ordered address: 0x78563412
```


inet_addr(), inet_aton ()

```
#include <arpa/inet.h>
in_addr_t inet_addr(const char *string);
int inet_aton(const char *string, struct in_addr *addr)
```

IPv4 address manipulation

- Convert dotted decimal IP address (String) to 32bit big-endian integer
- **inet_addr()**
 - Return value
 - Success: Internet address (32bit Integer)
 - Error: -1
- **inet_aton()**
 - Similar with **inet_addr()**, but the result value is returned using argument
 - Return value
 - Success: 1 (true)
 - Error: 0 (false)

■ Example

```
char *addr1 = "1.2.3.4";
char *addr2 = "127.232.124.79";
struct sockaddr_in addr_inet;

unsigned long conv_addr = inet_addr(addr1);
if (conv_addr == INADDR_NONE)
    printf("Error occurred! \n");
else
    printf("Network ordered integer addr#1: %#lx \n", conv_addr);

if (!inet_aton(addr, &addr_inet.sin_addr))
    printf("Conversion error\n");
else
    printf(" Network ordered integer addr#2: %#x \n", addr_inet.sin_addr.s_addr);

>> Network ordered integer addr: 0x4030201
>> Network ordered integer addr: 0x4f7ce87f
```

inet_ntoa()

```
#include <arpa/inet.h>
char *inet_ntoa(struct in_addr adr);
```

IPv4 address manipulation

- Convert 32bit big-endian integer to dotted decimal IP address (String)
- inet_ntoa()
 - Convert string IP address to dotted decimal IP address
 - Return value
 - Success: converted dotted decimal IP address (String)
 - Error: -1
 - Return value

■ Example

```
struct sockaddr_in addr1, addr2;
char *str_ptr;
char str_arr[20];

addr1.sin_addr.s_addr = htonl(0x1020304);
addr2.sin_addr.s_addr = htonl(0x1010101);

str_ptr = inet_ntoa(addr1.sin_addr);
strcpy(str_arr, str_ptr);
printf("Dotted-Decimal notation1: %s \n", str_ptr);
```

>> Dotted-Decimal notation1: 1.2.3.4

listen()

```
#include <sys/socket.h>
int listen(int sockfd, int backlog);
```

Listen for connections on a socket

- Tell OS to receive and queue SYN packets
- TCP Server only
- *sockfd*: file descriptor for the socket (socket type should be SOCK_STREAM)
- *backlog*: the maximum number of connection requests that system can queue while it waits for the server to accept them
- Return value
 - Success: 0
 - Error: -1
- Example

```
if (listen(serv_sock, 5) == -1)
    error_handling("listen() error");
```

accept()

```
#include <sys/socket.h>
```

```
int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);
```

Accept a connection on a socket

- TCP Server only
- Block until a connection request arrives
- *sockfd*: file descriptor for the accepted socket (socket type should be SOCK_STREAM)
- *addr*: pointer to a sockaddr structure to be filled in with the address of the client socket
- *addrlen*: it will contain the actual size of the client address

Return value

- Success: file descriptor for the accepted sock (>0)
- Error: -1

Example

```
clnt_addr_size = sizeof(clnt_addr);  
clnt_sock = accept(serv_sock, (struct sockaddr*)&clnt_addr, &clnt_addr_size);  
if (clnt_sock == -1)  
    error_handling("accept() error");
```

connect()

```
#include <sys/socket.h>
```

```
int connect(int sockfd, const struct sockaddr *addr, socklen_t addrlen);
```

Initiate a connection on a socket

- For a TCP socket, it establishes a connection to the server
- For a UDP socket, it simply stores the server's address so that the client can use a socket description
- *sockfd*: file descriptor for the socket
- *addr*: address to connect
- *addrlen*: the size (bytes) of the address structure pointed to by *addr*

- Return value
 - Success: 0
 - Failure: -1

■ Example

```
memset(&serv_addr, 0, sizeof(serv_addr));  
serv_addr.sin_family = AF_INET;  
serv_addr.sin_addr.s_addr = inet_addr(argv[1]);  
serv_addr.sin_port = htons(atoi(argv[2]));  
  
if (connect(sock, (struct sockaddr*)&serv_addr, sizeof(serv_addr)) == -1)  
    error_handling("connect() error!");
```

read()

```
#include <sys/socket.h>
ssize_t read(int fd, void *buf, size_t count);
```

Read from a file descriptor

- Block until data received
- Attempts to read up to *count* bytes from file descriptor *fd* (socket) into the buffer starting at *buf*
- Return value
 - Success: the number of bytes read is returned, and the file position is advanced by this number
 - Zero indicates end of file = connection close
 - Error: -1

■ Example

```
while (read_len = read(sock, &message[idx++], 1))
{
    if (read_len == -1)
    {
        error_handling("read() error!");
        break;
    }
    str_len += read_len;
}
```

write()

```
#include <sys/socket.h>
ssize_t write(int fd, const void *buf, size_t count);
```

Write to a file descriptor

- writes up to *count* bytes from the buffer starting at *buf* to the file (socket) referred to by the file descriptor *fd*
- Return value
 - Success: the number of bytes written is returned
 - Error: -1

■ Example

```
char message[]="Hello World!";
write(clnt_sock, message, sizeof(message));
```

recv()

```
#include <sys/socket.h>
```

```
ssize_t recv(int socket, void *buffer, size_t length, int flags);
```

Receive a message from a connected socket

- Block until data received
- *socket*: socket file descriptor
- *buffer*: points to a buffer where the message should be stored
- *length*: the length in bytes of the buffer pointed to by the *buffer* argument (maximum length of the *buffer*)
- *flags*: type of message reception
 - 0 for regular data
- Return value
 - Success: the number of bytes received
 - Zero indicates the connection close
 - Error: -1
- Example

```
while((str_len = recv(recv_sock, buf, sizeof(buf), 0)) != 0)
{
    if (str_len == -1)
        continue;
    buf[str_len]=0;
    puts(buf);
}
```


send()

```
#include <sys/socket.h>
ssize_t send(int socket, const void *buffer, size_t length, int flags);
```

Send a message on a socket

- Transmit the data in *buffer* upto *length* bytes
- *socket*: socket file descriptor
- *buffer*: buffer containing the message to send
- *length*: the length of the message in bytes.
- *flags*: type of message transmission
 - 0 for regular data
- Return value
 - Success: the number of bytes transmitted
 - Error: -1
- Example

```
write(sock, "123", strlen("123"));
send(sock, "4", strlen("4"), MSG_OOB);
```

close()

```
#include <unistd.h>
int close(int fd);
```

Close a file descriptor

- Prevent any more read and writes to the socket
- If the remote side calls `recv()`, it will return 0
- If the remote side calls `send()`, it will receive a signal SIGPIPE and `send()` will return -1 and `errno` will be set to EPIPE.
- Return value
 - Success: 0
 - Error: -1
- Example

```
close(clnt_sock);
close(serv_sock);
```

Review: tcp_server.c & tcp_client.c

// Server

```
serv_sock = socket(PF_INET, SOCK_STREAM, 0);  
if (serv_sock == -1)  
    error_handling("socket() error");
```

socket()

```
memset(&serv_addr, 0, sizeof(serv_addr));  
serv_addr.sin_family = AF_INET;  
serv_addr.sin_addr.s_addr = htonl(INADDR_ANY);  
serv_addr.sin_port = htons(atoi(argv[1]));
```

bind()

```
if (bind(serv_sock, (struct sockaddr*)&serv_addr, sizeof(serv_addr)) == -1)  
    error_handling("bind() error");
```

```
if (listen(serv_sock, 5) == -1)  
    error_handling("listen() error");
```

listen()

```
clnt_addr_size = sizeof(clnt_addr);  
clnt_sock = accept(serv_sock, (struct sockaddr*)&clnt_addr, &clnt_addr_size);  
if (clnt_sock == -1)  
    error_handling("accept() error");
```

accept()

```
write(clnt_sock, message, sizeof(message));
```

write()

```
close(clnt_sock);  
close(serv_sock);
```

close()

// Client

```
sock = socket(PF_INET, SOCK_STREAM, 0);  
if (sock == -1)  
    error_handling("socket() error");
```

socket()

```
memset(&serv_addr, 0, sizeof(serv_addr));  
serv_addr.sin_family = AF_INET;  
serv_addr.sin_addr.s_addr = inet_addr(argv[1]);  
serv_addr.sin_port = htons(atoi(argv[2]));
```

connect()

```
if (connect(sock, (struct sockaddr*)&serv_addr, sizeof(serv_addr)) == -1)  
    error_handling("connect() error!");
```

```
while (read_len = read(sock, &message[idx++], 1))  
{  
    if (read_len == -1)  
    {  
        error_handling("read() error!");  
        break;  
    }  
    str_len += read_len;  
}
```

read()

```
printf("Message from server: %s \n", message);  
printf("Function read call count: %d \n", str_len);
```

```
close(sock);
```

close()

Socket Programming: Roadmap

- Introduction
- TCP Socket programming #1
- **TCP Socket programming #2**
- UDP Socket programming
- Web Server Example

TCP: overview

RFCs: 793, 1122, 2018, 5681, 7323

- point-to-point:

- one sender, one receiver

- reliable, in-order *byte stream*:

- no “message boundaries”

- full duplex data:

- bi-directional data flow in same connection
- MSS: maximum segment size

- cumulative ACKs

- pipelining:

- TCP congestion and flow control set window size

- connection-oriented:

- handshaking (exchange of control messages) initializes sender, receiver state before data exchange

- flow controlled:

- sender will not overwhelm receiver

No Message boundary in TCP

- A "message boundary" is the separation between two messages being sent over a protocol.
- UDP preserves message boundaries
 - E.g., if the server calls write() twice, the client needs to call read() twice.
- TCP does not preserve message boundaries
 - E.g., even if the server calls write() twice, the client can read all data at once.

```
// Server (from tcp_server2.c)
char message1[] = "Hello World!111";
char message2[] = "Hello World!222";
...
write(clnt_sock, message1, sizeof(message1));
write(clnt_sock, message2, sizeof(message2));
```

```
// Client (from tcp_client2.c)
sleep(1);
read_len = read(sock, message, sizeof(message));
printf("Message from server: ");
for (i = 0; i < sizeof(message); i++)
    printf("%c", message[i]);
printf("Read_len: %d \n", read_len );
```

Results

```
yunmin@ym-ubuntu:~/Workspace$ ./hclient2 127.0.0.1 9191
Message from server: Hello World!111Hello World!222♦♦p^♦
Read_len: 32
```


Server Type

■ Iterative server

- A single server process receives and handles incoming requests on a “well-known” port
- Most UDP servers are iterative

■ Concurrent server

- A separate process to handle each client
- The main server process creates a new service process to handle each client
- Most TCP servers are concurrent: for ease of connection management



Iterative vs. Concurrent servers

Iterative server skeleton

```
int sockfd, newsockfd;
if ((sockfd = socket(...)) < 0)
    err_sys("socket error");
if ((bind(sockfd, ...) < 0)
    err_sys("bind error");
if (listen(sockfd, 5))
    err_sys("listen error");
for (;;) {
    newsockfd = accept(sockfd, ...);
    if (newsockfd < 0)
        err_sys("accept error");
    doit(newsockfd);
    close(newsockfd);
}
```

Concurrent server skeleton

```
int sockfd, newsockfd;
if ((sockfd = socket(...)) < 0)
    err_sys("socket error");
if ((bind(sockfd, ...) < 0)
    err_sys("bind error");
if (listen(sockfd, 5))
    err_sys("listen error");
for (;;) {
    newsockfd = accept(sockfd, ...);
    if (newsockfd < 0)
        err_sys("accept error");
    if (fork() == 0) {
        close(sockfd);
        doit(newsockfd);
        exit(0);
    }
    else {
        close(newsockfd);
    }
}
```


Echo Server Example: tcp_echo_server.c

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <arpa/inet.h>
#include <sys/socket.h>

#define BUF_SIZE 1024
void error_handling(char *message);

int main(int argc, char *argv[])
{
    int serv_sock, clnt_sock;
    char message[BUF_SIZE];
    int str_len, i;

    struct sockaddr_in serv_adr;
    struct sockaddr_in clnt_adr;
    socklen_t clnt_adr_sz;

    if (argc != 2) {
        printf("Usage : %s <port>\n", argv[0]);
        exit(1);
    }

    serv_sock=socket(PF_INET, SOCK_STREAM, 0);
    if (serv_sock == -1)
        error_handling("socket() error");
```

```
    memset(&serv_adr, 0, sizeof(serv_adr));
    serv_adr.sin_family = AF_INET;
    serv_adr.sin_addr.s_addr = htonl(INADDR_ANY);
    serv_adr.sin_port = htons(atoi(argv[1]));

    if (bind(serv_sock, (struct sockaddr*)&serv_adr, sizeof(serv_adr)) == -1)
        error_handling("bind() error");

    if (listen(serv_sock, 5) == -1)
        error_handling("listen() error");

    clnt_adr_sz=sizeof(clnt_adr);
    for (i = 0; i < 5; i++)
    {
        clnt_sock = accept(serv_sock, (struct sockaddr*)&clnt_adr, &clnt_adr_sz);
        if (clnt_sock == -1)
            error_handling("accept() error");
        else
            printf("Connected client %d \n", i+1);

        while ((str_len = read(clnt_sock, message, BUF_SIZE)) != 0)
            write(clnt_sock, message, str_len);

        close(clnt_sock);
    }

    close(serv_sock);
    return 0;
}
```

Echo Client Example: tcp_echo_client.c

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <arpa/inet.h>
#include <sys/socket.h>

#define BUF_SIZE 1024
void error_handling(char *message);

int main(int argc, char *argv[])
{
    int sock;
    char message[BUF_SIZE];
    int str_len;
    struct sockaddr_in serv_adr;

    if (argc != 3) {
        printf("Usage : %s <IP> <port>\n", argv[0]);
        exit(1);
    }

    sock=socket(PF_INET, SOCK_STREAM, 0);
    if (sock == -1)
        error_handling("socket() error");

    memset(&serv_adr, 0, sizeof(serv_adr));
    serv_adr.sin_family = AF_INET;
    serv_adr.sin_addr.s_addr = inet_addr(argv[1]);
    serv_adr.sin_port = htons(atoi(argv[2]));
```

```
    if (connect(sock, (struct sockaddr*)&serv_adr, sizeof(serv_adr)) == -1)
        error_handling("connect() error!");
    else
        puts("Connected.....");

    while(1)
    {
        fputs("Input message(Q to quit): ", stdout);
        fgets(message, BUF_SIZE, stdin);


        if (!strcmp(message, "q\n") || !strcmp(message, "Q\n"))
            break;

        str_len = write(sock, message, strlen(message));

        recv_len = 0;
        while (recv_len < str_len)
        {
            recv_cnt = read(sock, &message[recv_len], BUF_SIZE-1);
            if (recv_cnt == -1)
                error_handling("read() error!");
            recv_len+=recv_cnt;
        }
        message[recv_len]=0;
        printf("Message from server: %s", message);
    }
    close(sock);
    return 0;
}
```


Input & Output Buffer for TCP Socket

- Input buffer and output buffer exists on the TCP socket
- Input buffer and output buffer are created when socket is created
- Even if the socket is closed, data remaining in the output buffer will continue to be transmitted
- When the socket is closed, the data remaining in the input buffer is destroyed



Half-close

- Meaning of close()
 - It means complete destruction of the socket
 - No more I/O is possible
 - If the data transmitting and receiving have not been completed yet, it causes problems
- Half-close
 - Close only one of the input or output streams



shutdown()

```
#include <unistd.h>
int shutdown(int socket, int how);
```

Shut down socket send and receive operations

- This function shall cause all or part of a full-duplex connection on the socket
- *socket*: socket file descriptor
- *how*: type of shutdown
 - SHUT_RD (0): disables further receive operations
 - SHUT_WR (1): disables further send operations
 - SHUT_RDWR (2): disables further send and receive operations
- The shutdown() does not actually free up the socket descriptor. To free the descriptor, use close()

■ Return value

- Success: 0
- Error: -1

■ Example

```
shutdown(clnt_sd, SHUT_WR);
read(clnt_sd, buf, BUF_SIZE);
printf("Message from client: %s \n", buf);

close(clnt_sd); close(serv_sd);
```

File Server Example: tcp_file_server.c

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <arpa/inet.h>
#include <sys/socket.h>

#define BUF_SIZE 30
void error_handling(char *message);

int main(int argc, char *argv[])
{
    int serv_sd, clnt_sd;
    FILE * fp;
    char buf[BUF_SIZE];
    int read_cnt;

    struct sockaddr_in serv_adr, clnt_adr;
    socklen_t clnt_adr_sz;

    if (argc != 2) {
        printf("Usage: %s <port>\n", argv[0]);
        exit(1);
    }

    fp = fopen("file_server.c", "rb");
    serv_sd = socket(PF_INET, SOCK_STREAM, 0);
```

```
    memset(&serv_adr, 0, sizeof(serv_adr));
    serv_adr.sin_family = AF_INET;
    serv_adr.sin_addr.s_addr = htonl(INADDR_ANY);
    serv_adr.sin_port = htons(atoi(argv[1]));

    bind(serv_sd, (struct sockaddr*)&serv_adr, sizeof(serv_adr));
    listen(serv_sd, 5);

    clnt_adr_sz = sizeof(clnt_adr);
    clnt_sd = accept(serv_sd, (struct sockaddr*)&clnt_adr, &clnt_adr_sz);

    while(1)
    {
        read_cnt = fread((void*)buf, 1, BUF_SIZE, fp);
        if (read_cnt < BUF_SIZE)
        {
            write(clnt_sd, buf, read_cnt);
            break;
        }
        write(clnt_sd, buf, BUF_SIZE);
    }

    shutdown(clnt_sd, SHUT_WR);
    read(clnt_sd, buf, BUF_SIZE);
    printf("Message from client: %s\n", buf);

    fclose(fp);
    close(clnt_sd); close(serv_sd);
    return 0;
}
```

File Client Example: tcp_file_client.c

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <arpa/inet.h>
#include <sys/socket.h>

#define BUF_SIZE 30
void error_handling(char *message);

int main(int argc, char *argv[])
{
    int sd;
    FILE *fp;

    char buf[BUF_SIZE];
    int read_cnt;
    struct sockaddr_in serv_adr;
    if (argc != 3) {
        printf("Usage: %s <IP> <port>\n", argv[0]);
        exit(1);
    }
```

```
    fp = fopen("receive.dat", "wb");
    sd = socket(PF_INET, SOCK_STREAM, 0);

    memset(&serv_adr, 0, sizeof(serv_adr));
    serv_adr.sin_family = AF_INET;
    serv_adr.sin_addr.s_addr = inet_addr(argv[1]);
    serv_adr.sin_port = htons(atoi(argv[2]));

    connect(sd, (struct sockaddr*)&serv_adr, sizeof(serv_adr));


    while((read_cnt = read(sd, buf, BUF_SIZE)) != 0)
        fwrite((void*)buf, 1, read_cnt, fp);

    puts("Received file data");
    write(sd, "Thank you", 10);
    fclose(fp);
    close(sd);
    return 0;
}
```

Socket Programming: Roadmap

- Introduction
- TCP Socket programming #1
- TCP Socket programming #2
- **UDP Socket programming**
- Web Server Example

UDP Server/Client Function Call



recvfrom()

```
#include <sys/socket.h>
ssize_t recvfrom(int sockfd, void *buffer, size_t length, int flags,
                  struct sockaddr *src_addr, socklen_t *addrlen);
```

Receive a message from a socket

- Block until data received
- *sockfd*: socket file descriptor
- *buffer*: points to a buffer where the message should be stored
- *length*: receives the data up to *length* bytes into *buffer*
- *flags*: type of message reception
 - 0 for regular
- *src_addr*: pointer to a sockaddr structure to be filled in with the address of the peer socket
- *addrlen*: it will contain the actual size of the peer address
- Return value
 - Success: the number of bytes received
 - Zero for EOF
 - Error: -1
- Example

```
str_len = recvfrom(serv_sock, message, BUF_SIZE, 0,
                   (struct sockaddr*)&clnt_addr, &clnt_addr_sz);
```

sendto()

```
#include <sys/socket.h>
ssize_t sendto(int sockfd, const void *buffer, size_t length, int flags,
               const struct sockaddr *dest_addr, socklen_t *addrlen);
```

Send a message on a socket

- *sockfd*: socket file descriptor
- *buffer*: buffer containing the message to send
- *length*: transmits the data in *buffer* upto *length* bytes
- *flags*: type of message transmission
 - 0 for regular
- *dest_addr*: address to transmit data
- *addrlen*: the size (bytes) of the address structure pointed to by *addr*
- Return value
 - Success: the number of bytes sent
 - Error: -1
- Example

```
sendto(serv_sock, message, str_len, 0,
       (struct sockaddr*)&clnt_addr, clnt_addr_sz);
```

UDP Server Example: udp_echo_server.c

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <arpa/inet.h>
#include <sys/socket.h>

#define BUF_SIZE 30
void error_handling(char *message);

int main(int argc, char *argv[])
{
    int serv_sock;
    char message[BUF_SIZE];
    int str_len;
    socklen_t clnt_adr_sz;

    struct sockaddr_in serv_adr, clnt_adr;
    if (argc != 2) {
        printf("Usage : %s <port>\n", argv[0]);
        exit(1);
    }
```

```
serv_sock = socket(PF_INET, SOCK_DGRAM, 0);
if (serv_sock == -1)
    error_handling("UDP socket creation error");

memset(&serv_adr, 0, sizeof(serv_adr));
serv_adr.sin_family = AF_INET;
serv_adr.sin_addr.s_addr = htonl(INADDR_ANY);
serv_adr.sin_port = htons(atoi(argv[1]));

if (bind(serv_sock, (struct sockaddr*)&serv_adr, sizeof(serv_adr)) == -1)
    error_handling("bind() error");

while(1)
{
    clnt_adr_sz = sizeof(clnt_adr);
    str_len = recvfrom(serv_sock, message, BUF_SIZE, 0,
                      (struct sockaddr*)&clnt_adr, &clnt_adr_sz);
    sendto(serv_sock, message, str_len, 0,
           (struct sockaddr*)&clnt_adr, clnt_adr_sz);
}

close(serv_sock);
return 0;
}
```

UDP Client Example: udp_echo_client.c

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <arpa/inet.h>
#include <sys/socket.h>

#define BUF_SIZE 30
void error_handling(char *message);

int main(int argc, char *argv[])
{
    int sock;
    char message[BUF_SIZE];
    int str_len;
    socklen_t adr_sz;

    struct sockaddr_in serv_adr, from_adr;
    if (argc != 3) {
        printf("Usage : %s <IP> <port>\n", argv[0]);
        exit(1);
    }

    sock = socket(PF_INET, SOCK_DGRAM, 0);
    if (sock == -1)
        error_handling("socket() error");
```

socket()

```
memset(&serv_adr, 0, sizeof(serv_adr));
serv_adr.sin_family=AF_INET;
serv_adr.sin_addr.s_addr=inet_addr(argv[1]);
serv_adr.sin_port=htons(atoi(argv[2]));

while(1)
{
    fputs("Insert message(q to quit): ", stdout);
    fgets(message, sizeof(message), stdin);
    if (!strcmp(message, "q\n") || !strcmp(message, "Q\n"))
        break;

    sendto(sock, message, strlen(message), 0,
           (struct sockaddr*)&serv_adr, sizeof(serv_adr));
    adr_sz = sizeof(from_adr);
    str_len = recvfrom(sock, message, BUF_SIZE, 0,
                       (struct sockaddr*)&from_adr, &adr_sz);

    message[str_len]=0;
    printf("Message from server: %s", message);
}

close(sock);
return 0;
}
```

sendto()/
recvfrom()

close()

Connected UDP

- Internal procedures in sendto()
 - 1) Allocate a destination IP and port for UDP socket
 - 2) Send data to allocated destination IP and port
 - 3) Deallocate destination IP and port from UDP socket
- Connected UDP
 - We can use the socket that has destination address already assigned
 - Skip address allocate and deallocate procedure
 - We can use write() and receive() instead of sendto() and recvfrom()
 - It does not mean connection-oriented socket like TCP

How to create connected UDP

```
sock = socket(PF_INET, SOCK_DGRAM, 0);
if (sock == -1)
    error_handling("socket() error");

memset(&serv_adr, 0, sizeof(serv_adr));
serv_adr.sin_family = AF_INET;
serv_adr.sin_addr.s_addr = inet_addr(argv[1]);
serv_adr.sin_port = htons(atoi(argv[2]));

connect(sock, (struct sockaddr*)&serv_adr, sizeof(serv_adr));
```

UDP Client Example: udp_echo_con_client.c

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <arpa/inet.h>
#include <sys/socket.h>

#define BUF_SIZE 30
void error_handling(char *message);

int main(int argc, char *argv[])
{
    int sock;
    char message[BUF_SIZE];
    int str_len;
    socklen_t adr_sz;

    struct sockaddr_in serv_adr, from_adr;
    if (argc != 3) {
        printf("Usage : %s <IP> <port>\n", argv[0]);
        exit(1);
    }

    sock = socket(PF_INET, SOCK_DGRAM, 0);
    if (sock == -1)
        error_handling("socket() error");
```

```
    memset(&serv_adr, 0, sizeof(serv_adr));
    serv_adr.sin_family = AF_INET;
    serv_adr.sin_addr.s_addr = inet_addr(argv[1]);
    serv_adr.sin_port = htons(atoi(argv[2]));

    connect(sock, (struct sockaddr*)&serv_adr, sizeof(serv_adr));

    while(1)
    {
        fputs("Insert message(q to quit): ", stdout);
        fgets(message, sizeof(message), stdin);
        if (!strcmp(message, "q\n") || !strcmp(message, "Q\n"))
            break;

        write(sock, message, strlen(message));

        str_len = read(sock, message, sizeof(message)-1);
        message[str_len] = 0;

        printf("Message from server: %s", message);
    }
    close(sock);
    return 0;
}
```

Socket Programming: Roadmap

- Introduction
- TCP Socket programming #1
- TCP Socket programming #2
- UDP Socket Programming
- Web Server Example

Simple Web Server

- Let's implement a simple web server using TCP socket
- Our web server can provide simple web page by client request

Server

```
yunmin@ym-ubuntu:~/Workspace$ ./webserv_linux 9191  
Connection Request : 127.0.0.1:52208  
Connection Request : 127.0.0.1:52210
```

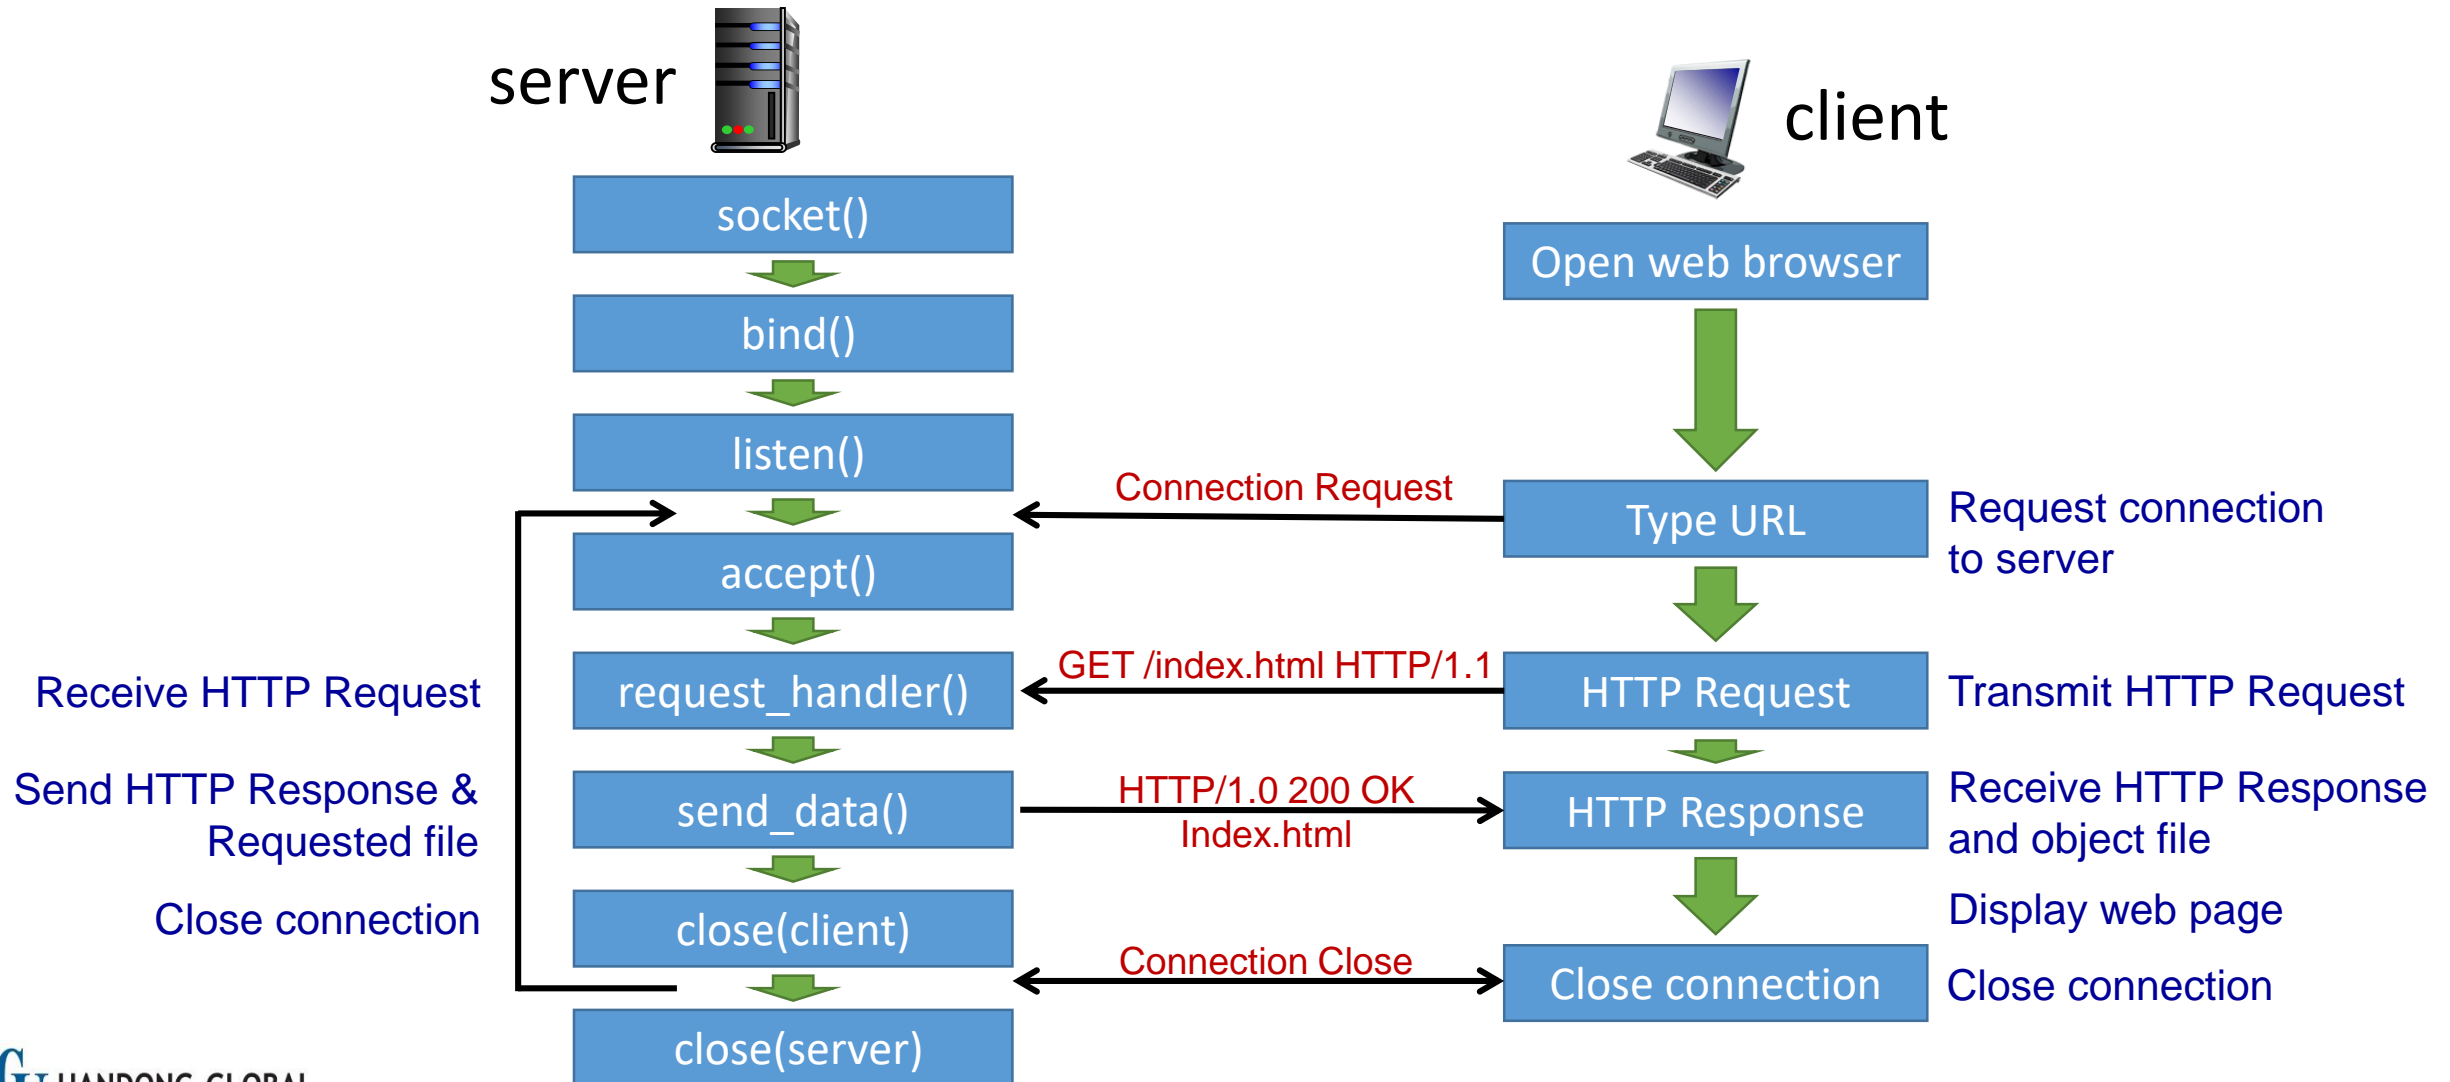
Client



Handong Global University!
Computer Network!
Socket Programming is Fun!

- We will implement iterative server, not concurrent server
- For the client, we will use web browser instead of implementing client program

Function Calls for Simple Web Server



Source Codes: main()

```
int main(int argc, char *argv[])
{
    int serv_sock, clnt_sock;
    struct sockaddr_in serv_adr, clnt_adr;
    int clnt_adr_size;
    char buf[BUF_SIZE];
    pthread_t t_id;
    if (argc != 2) {
        printf("Usage : %s <port>\n", argv[0]);
        exit(1);
    }

    serv_sock=socket(PF_INET, SOCK_STREAM, 0);
    memset(&serv_adr, 0, sizeof(serv_adr));
    serv_adr.sin_family = AF_INET;
    serv_adr.sin_addr.s_addr=htonl(INADDR_ANY);
    serv_adr.sin_port = htons(atoi(argv[1]));
    if (bind(serv_sock, (struct sockaddr*)&serv_adr, sizeof(serv_adr)) == -1)
        error_handling("bind() error");
    if (listen(serv_sock, 20) == -1)
        error_handling("listen() error");
```

```
while(1)
{
    clnt_adr_size = sizeof(clnt_adr);
    clnt_sock = accept(serv_sock, (struct sockaddr*)&clnt_adr, &clnt_adr_size);
    printf("Connection Request : %s:%d\n",
        inet_ntoa(clnt_adr.sin_addr), ntohs(clnt_adr.sin_port));

    request_handler(&clnt_sock);

    // pthread_create(&t_id, NULL, request_handler, &clnt_sock);
    // pthread_detach(t_id);
}
close(serv_sock);
return 0;
}
```

Source Codes: request_handler()

```
// void* request_handler(void *arg)
void request_handler(void *arg)
{
    int clnt_sock = *((int*)arg);
    char req_line[SMALL_BUF];
    FILE* clnt_read;
    FILE* clnt_write;

    char method[10];
    char ct[15];
    char file_name[30];

    clnt_read = fdopen(clnt_sock, "r");
    clnt_write = fdopen(dup(clnt_sock), "w");
    fgets(req_line, SMALL_BUF, clnt_read);
    if (strstr(req_line, "HTTP/") == NULL)
    {
        send_error(clnt_write);
        fclose(clnt_read);
        fclose(clnt_write);
        return;
    }
}
```

```
strcpy(method, strtok(req_line, " /"));
strcpy(file_name, strtok(NULL, " /"));
strcpy(ct, content_type(file_name));
if (strcmp(method, "GET")!=0)
{
    send_error(clnt_write);
    fclose(clnt_read);
    fclose(clnt_write);
    return;
}

fclose(clnt_read);
send_data(clnt_write, ct, file_name);
}
```

Source Codes: send_data()

```
void send_data(FILE* fp, char* ct, char* file_name)
{
    char protocol[] = "HTTP/1.0 200 OK\r\n";
    char server[] = "Server:Linux Web Server \r\n";
    char cnt_len[] = "Content-length:2048\r\n";
    char cnt_type[SMALL_BUF];
    char buf[BUF_SIZE];
    FILE* send_file;

    sprintf(cnt_type, "Content-type:%s\r\n\r\n", ct);
    send_file = fopen(file_name, "r");
    if (send_file == NULL)
    {
        send_error(fp);
        return;
    }
}
```

```
fputs(protocol, fp);
fputs(server, fp);
fputs(cnt_len, fp);
fputs(cnt_type, fp);

while (fgets(buf, BUF_SIZE, send_file) != NULL)
{
    fputs(buf, fp);
    fflush(fp);
}
fflush(fp);
fclose(fp);
}
```

Source Codes: content_type(), send_error()

```
char* content_type(char* file)
{
    char extension[SMALL_BUF];
    char file_name[SMALL_BUF];
    strcpy(file_name, file);
    strtok(file_name, ".");
    strcpy(extension, strtok(NULL, "."));

    if (!strcmp(extension, "html") || !strcmp(extension, "htm"))
        return "text/html";
    else
        return "text/plain";
}
```

```
void send_error(FILE* fp)
{
    char protocol[] = "HTTP/1.0 400 Bad Request\r\n";
    char server[] = "Server:Linux Web Server \r\n";
    char cnt_len[] = "Content-length:2048\r\n";
    char cnt_type[] = "Content-type:text/html\r\n\r\n";
    char content[] = "<html><head><title>NETWORK</title></head>"
                    "<body><font size=+5><br>Error! Please re-check the file name"
                    "and requet method"
                    "</font></body></html>";

    fputs(protocol, fp);
    fputs(server, fp);
    fputs(cnt_len, fp);
    fputs(cnt_type, fp);
    fflush(fp);
}
```

Appendix

Name-to-Address Conversion

gethostbyname()

```
#include <netdb.h>
struct hostent *gethostbyname(const char *hostname);
```

Get host information by taking host name


- We can obtain the host's official name, aliases, and IP addresses by *hostname*
- *hostname*: a hostname or an IPv4 address in standard dot notation
- Return value
 - Success: structure type of *hostent* for given host name
 - Error: NULL

struct hostent

- Defined in netdb.h

```
#define h_addr haddr_list[0]
```

```
struct hostent {  
    char    *h_name;           // official name of host  
    char    **h_aliases;       // alias list  
    int     h_addrtype;        // host address type  
    int     h_length;          // length of address  
    char    **h_addr_list;     // list of addresses  
};
```



gethostbyaddr()

```
#include <netdb.h>
struct hostent *gethostbyaddr(const char *addr, socklen_t len, int family);
```

Get host information by taking network byte order address

- We can obtain host's official name, aliases, and IP addresses by *addr* (network byte or address)
- *addr*: address of host (type = struct in_addr)
- *len*: length of *addr* (IPv4=4, IPv6=16)
- *family*: address family (AF_INET=IPv4, AF_INET6=IPv6)
- Return value
 - Success: structure type of *hostent* for given host name
 - Error: NULL